

Laboratory 3

(Due date: Oct. 10th)

OBJECTIVES

- Compile and execute multi-threaded C code in Ubuntu 12.04.4 using the Terasic DE2i-150 Development Kit.
- Learn multi-threading implementation using *pthread*s in C.
- Compare computation time of multi-threaded implementations using different number of threads.

REFERENCE MATERIAL

- Refer to the [board website](#) or the [Tutorial: Embedded Intel](#) for User Manuals and Guides.
- Refer to the [Tutorial: High-Performance Embedded Programming with the Intel® Atom™ platform](#) → *Tutorial 3 and 4* for associated examples.

ACTIVITIES

FIRST ACTIVITY: CENTERED MOVING AVERAGE (WINDOW SIZE = 7)

- Given an n -element vector \vec{a} , where $a(i)$ is an element of the vector ($i = 0, 1, \dots, n - 1$), the elements of the 7-element moving average \vec{f} are given by:

$$f(i) \leftarrow \frac{a(i-3) + a(i-2) + a(i-1) + a(i) + a(i+1) + a(i+2) + a(i+3)}{7}$$

- ✓ The moving average is usually a central moving average that can be computed using data equally spaced on either side of a central value (this needs the number of elements in the window to be odd).
- ✓ In the formula, $i = 0, 1, \dots, n - 1$. When the elements are not available (at the borders), we only use the available elements:

$$f(0) \leftarrow \frac{a(i) + a(i+1) + a(i+2) + a(i+3)}{4}$$

$$f(1) \leftarrow \frac{a(i-1) + a(i) + a(i+1) + a(i+2) + a(i+3)}{5}$$

$$f(2) \leftarrow \frac{a(i-2) + a(i-1) + a(i) + a(i+1) + a(i+2) + a(i+3)}{6}$$

$$f(n-1) \leftarrow \frac{a(i-3) + a(i-2) + a(i-1) + a(i)}{4}$$

$$f(n-2) \leftarrow \frac{a(i-3) + a(i-2) + a(i-1) + a(i) + a(i+1)}{5}$$

$$f(n-3) \leftarrow \frac{a(i-3) + a(i-2) + a(i-1) + a(i) + a(i+1) + a(i+2)}{6}$$

- Fig. 1 depicts an example. The original data (102 data points) is plotted as a series of dots. The 7-element moving average smooths short-term fluctuations and highlight longer-term trends.




Figure 1. Seven-element moving average

INSTRUCTIONS

- Write a .c program that reads in the parameter *nthreads*, reads the input data set from a binary input file (.bif), computes the 7-element centered moving average and displays the result.
 - ✓ Your code should measure the computation time (only the actual computation portion) in us.

Considerations:

- ✓ Input dataset: 100,000 elements of type `int32`. This is available in the provided `mydata.bif` file.
 - You can use this code snippet to read data from a binary file (use `typ=1` since each element is of type `int32`).

```
int read_binfile (int *data, int Length, char *in_file, int typ) {
    // data: array where the data read from file is placed
    // type: type = 0: each element is unsigned 8-bit integer. ==> 'unsigned char'
    //       type = 1: each element is a signed integer (32 bits) ==> 'int'
    // Length: # of elements to read (if type =1 --> number of 32-bit words)
    FILE *file_i;
    int i;
    size_t result, ELEM_SIZE;

    if (typ != 0 && typ != 1) { printf ("Wrong modifier (only 0, 1 accepted)\n"); return -1; }

    file_i = fopen(in_file,"rb");
    if (file_i == NULL) { printf ("Error opening file!\n"); return -1; }

    if (typ == 0) { // each element is an unsigned integer of 8 bits
        unsigned char *IM;
        IM = (unsigned char *) calloc (Length, sizeof(unsigned char));

        ELEM_SIZE = sizeof(unsigned char);
        result = fread (IM, sizeof(unsigned char), Length, file_i);

        for (i = 0; i < Length; i++) data[i] = (int) IM[i];
        free (IM); }
    else { // if (typ == 1) // each element is a signed 32-bit integer
        int *IM;
        IM = (int *) calloc (Length, sizeof(int));

        ELEM_SIZE = sizeof(int);
        result = fread (IM, sizeof(int), Length, file_i);

        for (i = 0; i < Length; i++) data[i] = IM[i];
        free (IM); }

    fclose (file_i);
    printf ("(read_binfile) Input binary file '%s': # of elements read = %ld\n", in_file, result);
    printf ("(read_binfile) Size of each element: %ld bytes\n", ELEM_SIZE);
    return 0;
}
```

- ✓ Strategy for parallelization: Given *nthreads* threads, the index *i* represents each thread from 0 to *nthreads*-1.
 - Each thread *i* is in charge of processing a slice of the input vector in order to generate a slice of the output vector.
 - The thread *i* computes the slice of the output vector \tilde{f} with the following indices:
 - From $\lfloor \frac{i \times n}{nthreads} \rfloor$ to $\lfloor \frac{(i+1) \times n}{nthreads} \rfloor - 1$.
 - Note that $nthreads \in [1, n]$.




Figure 2. Seven-element moving average for the 100,000-element input dataset.

- **Verification:** Fig. 2 depicts the input dataset along with the 7-element moving average.
 - ✓ The dataset is relatively large, so to verify the correctness of your result, have your program print out the following indices of output vector \vec{f} :
 - $f(0:19)$, $f(1000:1019)$, $f(99980:99999)$
 - ✓ Fig. 3 shows a screenshot of the execution in the Terminal with the three 20-element sets of values.

```

daniel@daniel-Inspiron-1545: ~/Dropbox/mystuff/work_ubuntu/labs/lab3
daniel@daniel-Inspiron-1545:~/Dropbox/mystuff/work_ubuntu/labs/lab3$ ./movavg_threads 5
(read_binfile) Input binary file 'mydata.bif': # of elements read = 100000
(read_binfile) Size of each element: 4 bytes

Creating 5 Threads
Thread 0 computes slice 0 (indices: 0 to 19999)
Thread 1 computes slice 1 (indices: 20000 to 39999)
Thread 2 computes slice 2 (indices: 40000 to 59999)
Thread 3 computes slice 3 (indices: 60000 to 79999)
Thread 4 computes slice 4 (indices: 80000 to 99999)
odata[0] = 218.7500    odata[1000] = 303.2857    odata[99980] = 397.7143
odata[1] = 273.6000    odata[1001] = 281.2857    odata[99981] = 406.1429
odata[2] = 310.8333    odata[1002] = 246.2857    odata[99982] = 384.2857
odata[3] = 329.2857    odata[1003] = 246.7143    odata[99983] = 405.2857
odata[4] = 277.7143    odata[1004] = 250.8571    odata[99984] = 381.5714
odata[5] = 244.8571    odata[1005] = 281.2857    odata[99985] = 320.4286
odata[6] = 310.8571    odata[1006] = 238.0000    odata[99986] = 287.5714
odata[7] = 311.0000    odata[1007] = 264.0000    odata[99987] = 245.7143
odata[8] = 252.1429    odata[1008] = 228.5714    odata[99988] = 194.7143
odata[9] = 220.1429    odata[1009] = 267.5714    odata[99989] = 179.5714
odata[10] = 216.8571   odata[1010] = 263.2857    odata[99990] = 178.7143
odata[11] = 210.7143   odata[1011] = 249.8571    odata[99991] = 181.0000
odata[12] = 238.1429   odata[1012] = 234.0000    odata[99992] = 234.4286
odata[13] = 235.8571   odata[1013] = 254.7143    odata[99993] = 262.2857
odata[14] = 272.8571   odata[1014] = 275.4286    odata[99994] = 319.8571
odata[15] = 300.0000   odata[1015] = 317.5714    odata[99995] = 374.8571
odata[16] = 302.2857   odata[1016] = 340.1429    odata[99996] = 376.7143
odata[17] = 252.2857   odata[1017] = 319.1429    odata[99997] = 371.5000
odata[18] = 316.5714   odata[1018] = 335.1429    odata[99998] = 391.0000
odata[19] = 335.0000   odata[1019] = 378.0000    odata[99999] = 377.2500
(write_binfile) Output binary file 'mydata.bof': # of (int32) elements written = 100000
start: 310084 us
end: 312837 us
Elapsed time (only actual computation): 2753 us
daniel@daniel-Inspiron-1545:~/Dropbox/mystuff/work_ubuntu/labs/lab3$

```

Figure 3. Execution of 7-element moving average showing three 20-element sets of values (the computation time corresponds to an execution on a Dell Inspiron laptop)

- Compile the code and execute the application on the DE2i-150 Board. Complete Table I (use an average of 10 executions in order to get the computation time for each case).
 - ✓ Example: `./my_movavg 10`
 - It will compute the moving average of the input dataset using 10 threads.

TABLE I. COMPUTATION TIME (US) VS. NUMBER OF THREADS

	<i>nthreads</i>									
	1	2	3	4	5	6	7	8	9	10
Computation Time (us)	19270.4	11164.8	10634.2	9907.4	11181.6	11497	11584.1	11034.9	12393	12187.8

- ✓ Comment on your results in Table I. Is there an optimal number of threads? At what point increasing the number of threads causes an increase in processing time?

The optimal number of threads is reached at $n = 4$. At this point the average computation time reaches its lowest time of 9907.7 us. After this point the average computation time reaches above 12000 us.

- Take a screenshot of the software running in the Terminal for $nthreads=5$. It should show the computation time along with the three 20-element sets of values for the output vector \vec{f} (like in Fig. 3).
- Provided file: `mydata.bif`.

SUBMISSION

- Demonstration: In this Lab 3, the requested screenshot of the software routine running in the Terminal suffices.
 - ✓ If you prefer, you can request a virtual session (Zoom) with the instructor and demo it.
- Submit to Moodle (an assignment will be created):
 - ✓ One .zip file
 - 1st Activity: The .zip file must contain the source files (`.c`, `.h`, `Makefile`), and the requested screenshot.
 - ✓ The lab sheet (a PDF file) with the completed Table I and your comments

TA signature: _____

Date: _____